

```

1  #!/usr/bin/env python3
2  # -*- coding: UTF-8 -*-
3
4  import argparse
5  import random
6  from contextlib import contextmanager
7  from multiprocessing import Pool, cpu_count
8  from timeit import default_timer as time
9
10 from tabulate import tabulate
11
12 CPU_COUNT = cpu_count()
13
14
15 class Timer:
16     """
17     Record timing information.
18     """
19
20     def __init__(self, *steps):
21         self._time_per_step = dict.fromkeys(steps)
22
23     def __getitem__(self, item):
24         return self._time_per_step[item]
25
26     @property
27     def time_per_step(self):
28         return {
29             step: elapsed_time
30             for step, elapsed_time in self._time_per_step.items()
31             if elapsed_time is not None and elapsed_time > 0
32         }
33
34     def start_for(self, step):
35         self._time_per_step[step] = -time()
36
37     def stop_for(self, step):
38         self._time_per_step[step] += time()
39
40
41 def merge_sort(array):
42     """Perform merge sort."""
43     array_length = len(array)
44
45     if array_length ≤ 1:
46         return array
47
48     middle_index = array_length // 2
49     left = merge_sort(array[:middle_index])
50     right = merge_sort(array[middle_index:])
51     return merge(left, right)
52
53
54 def merge(*arrays):
55     """Merge two sorted lists."""
56
57     # Support explicit left/right args, as well as a two-item
58     # tuple which works more cleanly with multiprocessing.
59     left, right = arrays[0] if len(arrays) == 1 else arrays

```

```

60     sorted_list = [0] * (len(left) + len(right))
61     i = j = k = 0
62
63     while i < len(left) and j < len(right):
64         if left[i] < right[j]:
65             sorted_list[k] = left[i]
66             i += 1
67         else:
68             sorted_list[k] = right[j]
69             j += 1
70         k += 1
71
72     while i < len(left):
73         sorted_list[k] = left[i]
74         i += 1
75         k += 1
76
77     while j < len(right):
78         sorted_list[k] = right[j]
79         j += 1
80         k += 1
81
82     return sorted_list
83
84
85 @contextmanager
86 def process_pool(size):
87     """Create a process pool and block until all processes have completed."""
88     pool = Pool(size)
89     yield pool
90     pool.close()
91     pool.join()
92
93
94 def parallel_merge_sort(array, ps_count):
95     """Perform parallel merge sort."""
96     timer = Timer("sort", "merge", "total")
97     timer.start_for("total")
98     timer.start_for("sort")
99
100     # Divide the list in chunks
101     step = int(len(array) / ps_count)
102
103     # Creates a pool of worker processes, one per CPU core.
104     # We then split the initial data into partitions, sized equally per
105     # worker, and perform a regular merge sort across each partition.
106     with process_pool(size=ps_count) as pool:
107         array = [array[i * step : (i + 1) * step] for i in range(ps_count)] + [array[ps_count * step
:]]
108
109         array = pool.map(merge_sort, array)
110
111         timer.stop_for("sort")
112         timer.start_for("merge")
113
114     # We can use multiprocessing again to merge sub-lists in parallel.
115     while len(array) > 1:
116         # If the number of partitions remaining is odd, we pop off the
117         # last one and append it back after one iteration of this loop,
118         # since we're only interested in pairs of partitions to merge.
119         extra = array.pop() if len(array) % 2 == 1 else None

```

```

119         array = [(array[i], array[i + 1]) for i in range(0, len(array), 2)]
120         array = pool.map(merge, array) + ([extra] if extra else [])
121
122         timer.stop_for("merge")
123         timer.stop_for("total")
124
125     final_sorted_list = array[0]
126
127     return timer, final_sorted_list
128
129
130 def get_command_line_parameters():
131     """Get the process count, array length from command line parameters."""
132
133     parser = argparse.ArgumentParser(
134         description="""Implement merge sort and multithreaded merge sort.
135         Compare the time required by both algorithms.
136         Also, analyze the performance of each algorithm for the best case and the worst case."""
137     )
138     parser.add_argument(
139         "-j",
140         "--jobs",
141         help="Number of processes to launch",
142         required=False,
143         default=CPU_COUNT,
144         type=lambda x: int(x)
145         if 0 < int(x) ≤ CPU_COUNT
146         else parser.error(f"Number of processes must be between 1 and {CPU_COUNT}"),
147     )
148     parser.add_argument(
149         "-l",
150         "--length",
151         help="Length of the array to sort",
152         required=False,
153         default=random.randint(3 * 10**6, 4 * 10**6), # Randomize the length of our list
154         type=lambda x: int(x) if 0 < int(x) else parser.error("Length of the array must be greater than
0"),
155     )
156     parser.add_argument(
157         "-a",
158         "--all",
159         help="Test all the variable length",
160         required=False,
161         default=False,
162         action="store_true",
163     )
164     return parser.parse_args()
165
166
167 def main(jobs, length, conclusion):
168     """Main function."""
169
170     main_timer = Timer("single_core", "list_generation")
171     main_timer.start_for("list_generation")
172
173     # Create an unsorted list with random numbers
174     randomized_array = [random.randint(0, i * 100) for i in range(length)]
175     main_timer.stop_for("list_generation")
176
177     print(f"List length: {length}")

```

```

178     print(f"Random list generated in {main_timer['list_generation']:.6f}s\n")
179
180     # Create a copy first due to mutation
181     randomized_array_sorted = randomized_array[:]
182     randomized_array_sorted.sort()
183
184     # Start timing the single-core procedure
185     print("Starting simple sort.")
186     main_timer.start_for("single_core")
187     sorted_array = merge_sort(randomized_array)
188     main_timer.stop_for("single_core")
189
190     # Comparison with Python list sort method
191     # serves also as validation of our implementation.
192     assert sorted_array == randomized_array_sorted, "The sorted array is not correct."
193     print(f"Single Core elapsed time: {main_timer['single_core']:.6f}s\n")
194
195     print("Starting parallel sort.")
196     parallel_timer, parallel_sorted_array = parallel_merge_sort(randomized_array, jobs)
197     print(f"Final merge duration: {parallel_timer['merge']:.6f}s")
198
199     assert parallel_sorted_array == randomized_array_sorted, "The sorted array is not correct."
200     print(f"{jobs}-Core elapsed time: {parallel_timer['total']:.6f}s\n" + "-" * 40, "\n")
201
202     conclusion.append([length, main_timer["single_core"], parallel_timer["total"]])
203
204
205 if __name__ == "__main__":
206     parameters = get_command_line_parameters()
207
208     jobs = parameters.jobs
209     length = parameters.length
210     all_cases = parameters.all
211     conclusion = []
212     print(f"Using {jobs} cores\n")
213
214     if all_cases:
215         l = 1
216         while l < 10**8:
217             main(jobs, l, conclusion)
218             l *= 10
219     else:
220         main(jobs, length, conclusion)
221
222     print(tabulate(conclusion, headers=["Array Length", "Single-Threaded", "Multi-Threaded"],
223         tablefmt="outline"))
224
225 """
226 OUTPUT:> python3 merge_sort.py -j 16
227
228 Using 16 cores
229
230 List length: 3332517
231 Random list generated in 1.826744s
232
233 Starting simple sort.
234 Single Core elapsed time: 14.580756s
235
236 Starting parallel sort.
237 Final merge duration: 2.910911s

```

```

237 16-Core elapsed time: 5.131196s
238 -----
239
240 +-----+-----+-----+
241 |   Array Length |   Single-Threaded |   Multi-Threaded |
242 +-----+-----+-----+
243 |       3332517 |       14.5808 |       5.1312 |
244 +-----+-----+-----+
245 ""
246
247 ""
248 OUTPUT:> python3 merge_sort.py -a
249
250 Using 16 cores
251
252 List length: 1
253 Random list generated in 0.000005s
254
255 Starting simple sort.
256 Single Core elapsed time: 0.000003s
257
258 Starting parallel sort.
259 Final merge duration: 0.003702s
260 16-Core elapsed time: 0.028538s
261 -----
262
263 List length: 10
264 Random list generated in 0.000035s
265
266 Starting simple sort.
267 Single Core elapsed time: 0.000023s
268
269 Starting parallel sort.
270 Final merge duration: 0.002112s
271 16-Core elapsed time: 0.024624s
272 -----
273
274 List length: 100
275 Random list generated in 0.000125s
276
277 Starting simple sort.
278 Single Core elapsed time: 0.000198s
279
280 Starting parallel sort.
281 Final merge duration: 0.001935s
282 16-Core elapsed time: 0.024828s
283 -----
284
285 List length: 1000
286 Random list generated in 0.000616s
287
288 Starting simple sort.
289 Single Core elapsed time: 0.002207s
290
291 Starting parallel sort.
292 Final merge duration: 0.004501s
293 16-Core elapsed time: 0.042195s
294 -----
295
296 List length: 10000

```

```

297 Random list generated in 0.006530s
298
299 Starting simple sort.
300 Single Core elapsed time: 0.024342s
301
302 Starting parallel sort.
303 Final merge duration: 0.009236s
304 16-Core elapsed time: 0.035944s
305 -----
306
307 List length: 100000
308 Random list generated in 0.061283s
309
310 Starting simple sort.
311 Single Core elapsed time: 0.297823s
312
313 Starting parallel sort.
314 Final merge duration: 0.081855s
315 16-Core elapsed time: 0.166326s
316 -----
317
318 List length: 1000000
319 Random list generated in 0.612833s
320
321 Starting simple sort.
322 Single Core elapsed time: 3.685065s
323
324 Starting parallel sort.
325 Final merge duration: 0.645426s
326 16-Core elapsed time: 1.235625s
327 -----
328
329 List length: 10000000
330 Random list generated in 5.445338s
331
332 Starting simple sort.
333 Single Core elapsed time: 45.568661s
334
335 Starting parallel sort.
336 Final merge duration: 5.418382s
337 16-Core elapsed time: 11.996948s
338 -----
339
340 +-----+-----+-----+
341 | Array Length | Single-Threaded | Multi-Threaded |
342 +-----+-----+-----+
343 |          1 |      2.934e-06 |      0.0285377 |
344 |         10 |      2.2908e-05 |      0.0246235 |
345 |        100 |    0.000198492 |      0.0248281 |
346 |       1000 |    0.00220744 |      0.0421945 |
347 |      10000 |    0.0243417 |      0.0359442 |
348 |     100000 |    0.297823 |      0.166326 |
349 |    1000000 |    3.68507 |      1.23563 |
350 |   10000000 |   45.5687 |     11.9969 |
351 +-----+-----+-----+
352 ""
353

```